

Table of Contents

<i>Part 1: Design Document</i>	2
A1: Business Problem	2
A2: Proposed Data Structure	2
A3: Database Justification	3
A4: Data Utilization	4
B: Logical Data Model	5
C: Database Objects	5
D: Scalability Strategies	5
E: Privacy & Security Measures	6
<i>Part 2: Implementation</i>	7
F1 – Write script to create a database instance named “D597 Task 1” using SQL	7
F2 – Write script to import the data records from the Scenario 2 CSV file into the D597 database	8
F3 – Write script for three queries to retrieve specific information from EcoMart database.	11
F4 - Apply optimization techniques to improve the run time of your queries from part F3, providing output results via a screenshot.	14
<i>Part 3: Presentation</i>	21

Part 1: Design Document

A1: Business Problem

EcoMart faces challenges managing diverse sustainability attributes across their product offering from eco-friendly vendors. This limits their commitment to providing, promoting, and connecting sustainability and environmental consciousness to consumers.

EcoMart is looking to create a database to support their business goals and values by addressing these problems:

- Lacking platform scalability and flexibility for their consumers
- Inconsistent tracking of sustainability certifications
- Difficulty linking eco-friendly products to verification of sustainability certifications
- Sacrificing platform speed and reliability for consumers

A2: Proposed Data Structure

To solve these problems, EcoMart will adopt a **relational database** that highlights the following tables, entities, and attributes:

Product Table - *Products*

- ProductID (Integer) - PK
- ProductName (Varchar)
- ProdDescription (Varchar)
- Price (Decimal 10,2)
- CategoryID (Integer) – FK
- BrandID (Integer) - FK
- CategoryName (Varchar)

Customer Table - *Customers*

- CustomerID (Integer) - PK
- CustomerName (Varchar)
- CustomerPhone (Varchar)
- CustomerEmail (Varchar)

Categories Table – *Categories*

- CategoryID (Integer) – PK
- CategoryName (Varchar)

Brands Table – *Brands*

- BrandID (Integer) - PK
- BrandName (Varchar)

Orders Table – *Orders*

- OrderID (Integer) - PK
- CustomerID (Integer) - FK

- OrderDate (Date)
- OrderStatus (Varchar)
- TotalAmount (Decimal 10,2)

Order Details Table – *OrderDetails*

- OrderDetailID (Integer) - PK
- OrderID (Integer) - FK
- ProductID (Integer) - FK
- Quantity (Integer)
- UnitPrice (Decimal 10,2)

Reviews Table – *Reviews*

- ReviewID (Integer) - PK
- CustomerID (Integer) - FK
- ProductID (Integer) - FK
- Rating (TinyInteger)
- ReviewText (Varchar)
- ReviewDate (Date)

Certifications Table – *Certifications*

- CertificationID (Integer) - PK
- CertificationName (Varchar)

Inventory Table – *Inventory*

- ProductID (Integer) – PK, FK
- StockLevel (Integer)
- ReorderLevel (Integer)

Product Certifications Table - *ProductCertifications*

- ProductID (Integer) - FK
- CertificationID (Integer) - FK

Supplier Table - *Suppliers*

- SupplierID (Integer) – PK
- ProductID (Integer) - FK
- SupplierName (Varchar)
- ProdAvailability (Boolean)
- SustainabilityCert (Boolean)

A3: Database Justification

A relational database for EcoMart solves these issues by:

1. Supporting **complex relationships** between entities such as, products, certifications, customers, suppliers, etc.

2. Enforcing **ACID compliance (atomicity, consistency, isolation, durability)** and **data integrity** for accurate inventory/sustainability updates
3. Enabling **OLAP (online analytical processing) queries** for flexibility and high-quality data reporting
4. Providing **scalability and performance optimization** for past, current, and new orders, products and customers.

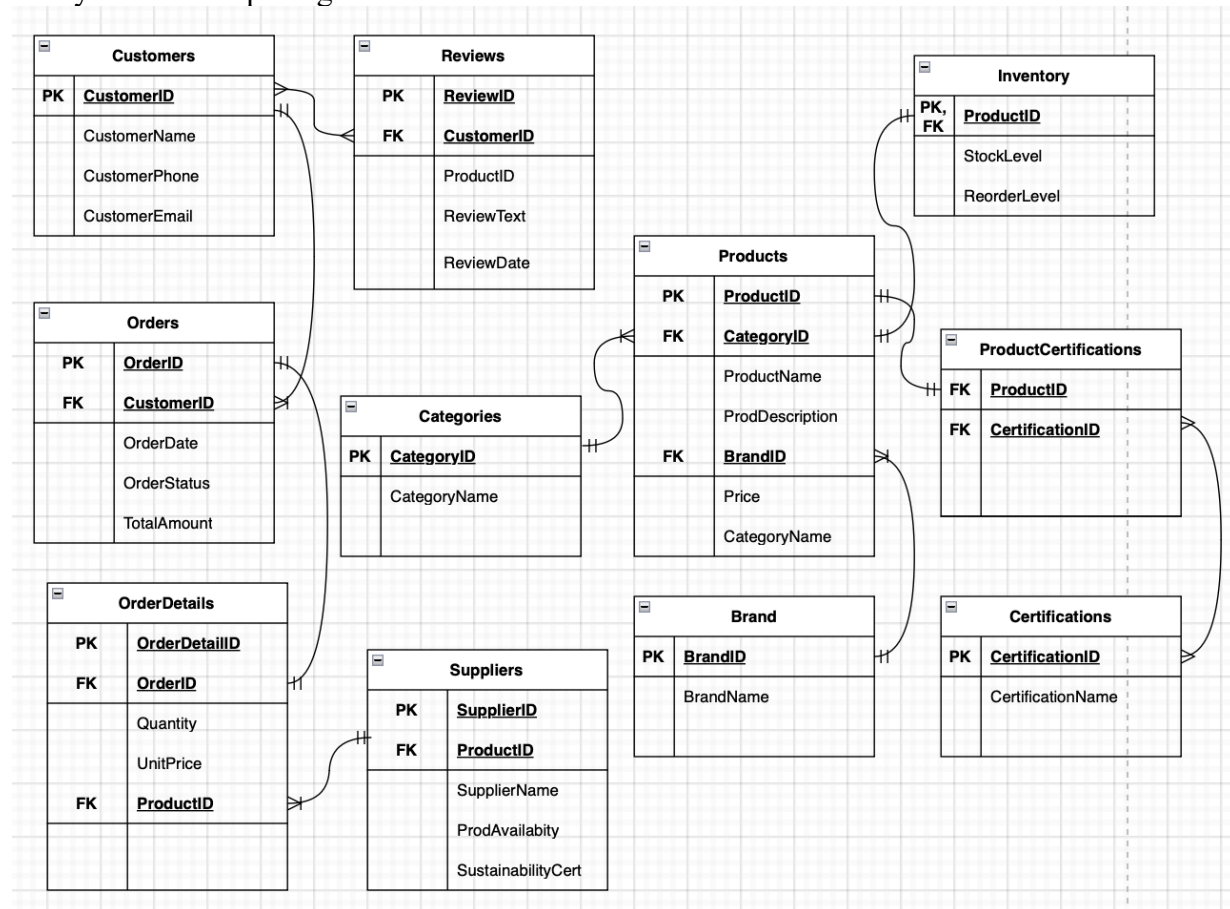
A4: Data Utilization

The data utilization for EcoMart is as follows:

- **Product discovery** – the ability for customers to search and filter through products using entities such as, categorization, certification, availability, and price.
- **Order processing** – the system can track customer orders, inventory updates, and fulfillment details.
- **Customer engagement** – the process of submitting reviews and ratings for product support
- **Sustainability analytics** – the ability to analyze eco-friendly trends and certification statuses

B: Logical Data Model

Entity-Relationship Diagram for EcoMart



C: Database Objects

The database for EcoMart will contain the following:

- **Tables** – this database object is how rows and columns will be stored
- **Indexes** – this database object is used to optimize queries so the performance is enhanced
- **Views** – this database object is for simplifying complex queries to support data presentation

The files within the database will store attributes such as data types, default values, constraints, and if an attribute is NULL or NOT NULL.

D: Scalability Strategies

To support scalability and future growth, we will use the following strategies:

- **Denormalization** – Combining and adding necessary information amongst tables such as creating CategoryName and SupplierName directly into the Products table.
- **Horizontal Scaling (Sharding)** – Dividing the order and supplier data into geographic regions or time zones to support high volumes of orders and shipping.

- **Vertical Scaling** – Increasing resources to servers such as expanding storage to be able to handle more orders in the future while keeping a log of past and current orders in the database.
- **Caching** – Implementing caching for specific “every day” queries such as finding most sold products or products with certifications. This will save time spent on rewriting popular queries.

E: Privacy & Security Measures

The privacy and security measures that should be implemented are as follows:

- 1. Data Protection (Encryption):**
 - a. Sensitive data like customer information (contact info, billing, etc.) needs to be encrypted and protected from possible data breaches.
- 2. Access Control:**
 - a. Assignment of roles such as a “sustainability auditor” should be implemented to restrict data visibility on specific and relevant tables
- 3. Compliance:**
 - a. There are laws that deal with storing and handling personal information and data. There needs to be compliance checks in place to ensure these laws (GDPR, PCI-DSS, etc.) are being upheld.
- 4. Monitoring:**
 - a. The monitoring of data is important because as EcoMart grows, it’s more susceptible to cyber-attacks and threats
 - b. There should be an implementation for real-time alerting for suspicious queries, such as bulk supplier data exports.

Part 2: Implementation

F1 – Write script to create a database instance named “D597 Task 1” using SQL

Create - Database [X]

General Definition Security Parameters Advanced SQL

```
1 CREATE DATABASE "D597 Task 1"
2 WITH
3 OWNER = postgres
4 ENCODING = 'UTF8'
5 LOCALE_PROVIDER = 'libc'
6 CONNECTION LIMIT = -1
7 IS_TEMPLATE = False;
```

[i] [?] [X Close] [Reset] [Save]

▼ Servers

- > PostgreSQL 17.4
- ▼ PostgreSQL 17
 - ▼ Databases (5)
 - ▼ **D597 Task 1**
 - > Casts
 - > Catalogs
 - > Event Triggers
 - > Extensions
 - > Foreign Data Wrappers
 - > Languages
 - > Publications
 - ▼ Schemas (1)
 - > public
 - > Subscriptions

F2 – Write script to import the data records from the Scenario 2 CSV file into the D597 database

Query	Query History
44	--CREATE TABLE in D597 Task1 Database
45	CREATE TABLE EcoMart (
46	region VARCHAR(100),
47	country VARCHAR(100),
48	item_Type VARCHAR(100),
49	sales_Channel VARCHAR(50),
50	order_Priority VARCHAR(1),
51	order_Date DATE,
52	order_ID INT,
53	ship_Date DATE,
54	units_Sold INT,
55	unit_Price DECIMAL(10,2),
56	unit_Cost DECIMAL(10,2),
57	total_Revenue DECIMAL(10,2),
58	total_Cost DECIMAL(10,2),
59	total_Profit DECIMAL(10,2)
60)
61	

Data Output	Messages	Notifications
CREATE TABLE		
Query returned successfully in 60 msec.		

COPY data from CSV file

Query	Query History
61	
62	--COPY command to copy the data in EcoMart table
63	COPY ecomart
64	FROM '/Applications/PostgreSQL 17/100000_Sales_Records.csv'
65	DELIMITER ','
66	CSV HEADER;
67	

Data Output	Messages	Notifications
COPY 100000		
Query returned successfully in 404 msec.		

Query Query History

```

61
62 --COPY command to copy the data in EcoMart table
63 COPY ecomart
64 FROM '/Applications/PostgreSQL 17/100000_Sales_Records.csv'
65 DELIMITER ','
66 CSV HEADER;
67
68 --COUNT to ensure amount of records in dataset (100,000)
69 SELECT COUNT(*)
70 FROM ecomart;
71

```

Data Output Messages Notifications

Showing row

	count bigint
1	100000

SELECT the entire dataset for EcoMart

D597 Task 1/postgres@PostgreSQL 17

Query Query History Scratch Pad

```

1 SELECT * FROM ecomart

```

Data Output Messages Notifications

Showing rows: 1 to 1000 Page No: 1 of 100

	region character varying (100)	country character varying (100)	item_type character varying (100)	sales_channel character varying (50)	order_priority character varying (1)
1	Middle East and North Africa	Azerbaijan	Snacks	Online	C
2	Central America and the Caribbean	Panama	Cosmetics	Offline	L
3	Sub-Saharan Africa	Sao Tome and Principe	Fruits	Offline	M
4	Sub-Saharan Africa	Sao Tome and Principe	Personal Care	Online	M
5	Central America and the Caribbean	Belize	Household	Offline	H
6	Europe	Denmark	Clothes	Online	C
7	Europe	Germany	Cosmetics	Offline	M
8	Middle East and North Africa	Turkey	Fruits	Online	C
9	Europe	United Kingdom	Snacks	Online	H
10	Asia	Kazakhstan	Cosmetics	Online	H
11	Central America and the Caribbean	Haiti	Cosmetics	Online	C
12	Europe	Italy	Clothes	Online	M
13	Europe	Malta	Household	Offline	L
14	Middle East and North Africa	Jordan	Household	Offline	L
15	Asia	Cambodia	Vegetables	Offline	H
16	Central America and the Caribbean	Saint Kitts and Nevis	Office Supplies	Online	H
17	Sub-Saharan Africa	Cameroon	Fruits	Online	H
18	Middle East and North Africa	Bahrain	Vegetables	Offline	L

Total rows: 100000 Query complete 00:00:00.191 LF Ln 1, Col 22

INSERT INTO command to insert new data entry (new order)

Query

Query History

90

91

92

93

94

95

96

97

98

99

100

101

102

--INSERT INTO command to insert new order into EcoMart table

INSERT INTO ecomart(region, country, item_type, sales_channel,

order_priority, order_date, order_id, ship_date, units_sold,

unit_price, unit_cost, total_revenue, total_cost, total_profit)

VALUES

('North America', 'Canada',

'Beverages', 'Online', 'H', '5/27/25',

'999999998', '6/3/25', '8759', '47.45',

'31.79', '415614.55', '278448.61', '137165.94');

SELECT * FROM ecomart WHERE order_id = '999999998';

Data Output

Messages

Notifications

INSERT 0 1

Query returned successfully in 75 msec.

Query

Query History

90

91

92

93

94

95

96

97

98

99

100

101

102

103

--INSERT INTO command to insert new order into EcoMart table

INSERT INTO ecomart(region, country, item_type, sales_channel,

order_priority, order_date, order_id, ship_date, units_sold,

unit_price, unit_cost, total_revenue, total_cost, total_profit)

VALUES

('North America', 'Canada',

'Beverages', 'Online', 'H', '5/27/25',

'999999998', '6/3/25', '8759', '47.45',

'31.79', '415614.55', '278448.61', '137165.94');

SELECT region, country, item_type, order_id

FROM ecomart WHERE order_id = '999999998';

Data Output

Messages

Notifications

≡+

📄

▼

📋

▼

🗑

🗑

📥

⬇

📈

SQL

Showing rows: 1 to

	region character varying (100) 🔒	country character varying (100) 🔒	item_type character varying (100) 🔒	order_id integer 🔒
1	North America	Canada	Beverages	999999998











F3 – Write script for three queries to retrieve specific information from EcoMart database.

Q1 – What are the top 5 products by total profit?



[Query](#) [Query History](#)

```
1  --Q1 - What are the top 5 products by total profit?
2  ✓ SELECT
3      item_type,
4      SUM (total_profit) as total_profit
5  FROM ecomart
6  GROUP BY item_type
7  ORDER BY total_profit DESC
8  LIMIT 5;
```

[Data Output](#) [Messages](#) [Notifications](#)



Showing row

	item_type character varying (100) 	total_profit numeric 
1	Cosmetics	7289406555.68
2	Household	6870966095.35
3	Office Supplies	5339532912.50
4	Baby Food	4017647893.20
5	Cereal	3743318890.62

[Query](#) [Query History](#)

```
1  --Q1 - What are the top 5 products by total profit?
2  ✓ SELECT
3      item_type,
4      SUM (total_profit) as total_profit
5  FROM ecomart
6  GROUP BY item_type
7  ORDER BY total_profit DESC
8  LIMIT 5;
```

[Data Output](#) [Messages](#) [Notifications](#)







Successfully run. Total query runtime: 131 msec.
5 rows affected.

Q2 – What is the average delivery time by sales channel?

[Query](#) [Query History](#)

```
1  --Q2 - What is the average delivery time by sales channel?
2  SELECT
3      sales_channel,
4      AVG(ship_date - order_date) AS avg_delivery_days
5  FROM ecomart
6  GROUP BY sales_channel;
```

[Data Output](#) [Messages](#) [Notifications](#)



Showing rows: 1 to 2

	sales_channel character varying (50)	avg_delivery_days numeric
1	Offline	24.9830817282665279
2	Online	25.0886642426179726

[Query](#) [Query History](#)

```
10  --Q2 - What is the average delivery time by sales channel?
11  SELECT
12      sales_channel,
13      AVG(ship_date - order_date) AS avg_delivery_days
14  FROM ecomart
15  GROUP BY sales_channel;
```

[Data Output](#) [Messages](#) [Notifications](#)

Successfully run. Total query runtime: 94 msec.
2 rows affected.

Q3 – What is the total revenue by region?

Query

Query History

1

--Q3 - What is the total revenue by region?

2

▼

SELECT

3

region,

4

SUM(total_revenue) AS total_revenue

5

FROM ecomart

6

GROUP BY region

7

ORDER BY total_revenue DESC;

8

Data Output

Messages

Notifications

≡+

▼

▼

SQL

	region character varying (100)	total_revenue numeric
1	Sub-Saharan Africa	34958453406.17
2	Europe	34241150923.39
3	Asia	19293401219.82
4	Middle East and North Africa	16921412794.52
5	Central America and the Caribbean	14553730165.29
6	Australia and Oceania	10701522223.73
7	North America	2937002333.49

17

--Q3 - What is the total revenue by region?

18

▼

SELECT

19

region,

20

SUM(total_revenue) AS total_revenue

21

FROM ecomart

22

GROUP BY region

23

ORDER BY total_revenue DESC;

24

Data Output

Messages

Notifications

Successfully run. Total query runtime: 77 msec.
7 rows affected.

F4 - Apply optimization techniques to improve the run time of your queries from part F3, providing output results via a screenshot.

1. Creating indexes for the frequently used columns (item types, regions, order dates, sales channel)

Query Query History

```
1 --F4
2 --Create Indexes for frequently used columns
3 CREATE INDEX idx_item_type ON ecomart(item_type);
4 CREATE INDEX idx_region ON ecomart(region);
5 CREATE INDEX idx_order_date ON ecomart(order_date);
6 CREATE INDEX idx_sales_channel ON ecomart(sales_channel);
7
8 SELECT * FROM pg_indexes WHERE tablename = 'ecomart';
9
```

Data Output Messages Notifications

CREATE INDEX

Query returned successfully in 243 msec.

D597 Task 1/postgres@Postgre SQL 17.4

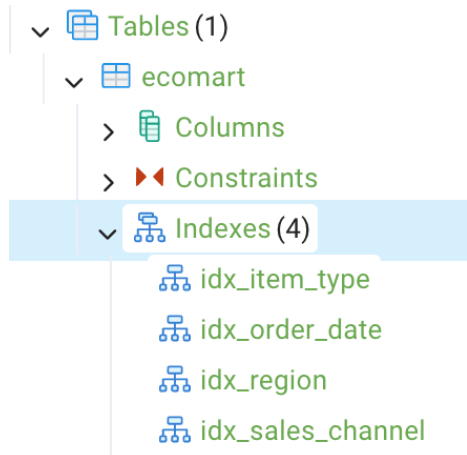
Query Query History Scratch Pad X

```
1 CREATE INDEX idx_item_type ON ecomart(item_type);
2 CREATE INDEX idx_region ON ecomart(region);
3 CREATE INDEX idx_order_date ON ecomart(order_date);
4 CREATE INDEX idx_sales_channel ON ecomart(sales_channel);
5
6 SELECT * FROM pg_indexes WHERE tablename = 'ecomart';
```

Data Output Messages Notifications

Showing rows: 1 to 4 Page No: 1 of 1

	schemaname name	tablename name	indexname name	tablespace name	indexdef text
1	public	ecomart	idx_item_type	[null]	CREATE INDEX idx_item_type ON public.ecomart USING btree (item_type)
2	public	ecomart	idx_region	[null]	CREATE INDEX idx_region ON public.ecomart USING btree (region)
3	public	ecomart	idx_order_date	[null]	CREATE INDEX idx_order_date ON public.ecomart USING btree (order_date)
4	public	ecomart	idx_sales_channel	[null]	CREATE INDEX idx_sales_channel ON public.ecomart USING btree (sales_channel)



2. Creating VIEWS to store summary tables for Total Revenue by Region, Top Profit Products, Avg Delivery Time by Sales Channel

Q1- Create a view for Total Revenue by Region

D597 Task 1/postgres@Postgre SQL 17.4

Query Query History

```

1  --Q1: Create a view for Total Revenue by Region
2  CREATE VIEW total_revenue_region AS
3  SELECT
4      region,
5      SUM(total_profit) AS total_profit
6  FROM ecomart
7  GROUP BY region;
8
9  --Output: The created VIEW for Total Revenue by Region
10 SELECT * FROM total_revenue_region;
11

```

Data Output Messages Notifications

Showing rows:

	region character varying (100)	total_profit numeric
1	Australia and Oceania	3175423561.38
2	Asia	5707511516.76
3	Central America and the Caribbean	4287210522.47
4	North America	872551616.84
5	Sub-Saharan Africa	10306312642.23
6	Middle East and North Africa	4979534378.88
7	Europe	10080579491.05

Before Optimization (Q1 - Total Revenue by Region):

```
23 --Before Optimization / No indexes or views applied yet
24 ✓ SELECT region, SUM(total_revenue) AS total_revenue
25 FROM ecomart
26 GROUP BY region
27 ORDER BY total_revenue DESC;
28
```

Data Output Messages Notifications

Successfully run. Total query runtime: 90 msec.
7 rows affected.



After Optimization (Q1 - Total Revenue by Region):








```
20 --Output: The created VIEW for Total Revenue by Region
21 SELECT * FROM total_revenue_region;
22
```

Data Output Messages Notifications

Successfully run. Total query runtime: 85 msec.
7 rows affected.

Q2 – Create a view for Top Profit Products

 D597 Task 1/postgres@Postgre SQL 17.4 

    No limit   

Query







Query History



```
14 --Q2: Create a view for Top Profit Products
15 CREATE VIEW top_profit_products AS
16 SELECT
17     item_type,
18     SUM(total_profit) AS total_profit
19 FROM ecomart
20 GROUP BY item_type;
21
22 --Output: The created VIEW for Top Profit Products
23 SELECT * FROM top_profit_products;
```

Data Output

Messages

Notifications

      SQL Showing rows

	item_type character varying (100) 	total_profit numeric 
1	Fruits	98321435.16
2	Baby Food	4017647893.20
3	Beverages	650112575.58
4	Vegetables	2604417796.68
5	Clothes	3067841433.60
6	Cereal	3743318890.62
7	Cosmetics	7289406555.68
8	Meat	2387834992.40
9	Office Supplies	5339532912.50
10	Household	6870966095.35
11	Snacks	2299287932.88
12	Personal Care	1040435215.96

Before Optimization (Q2 - Top Profit Products):

```
43 --Before Optimization / No indexes or views applied yet
44 ✓ SELECT item_type, SUM(total_profit) AS total_profit
45 FROM ecomart
46 GROUP BY item_type
47 ORDER BY total_profit DESC
```

Data Output Messages Notifications

Successfully run. Total query runtime: 113 msec.
12 rows affected.

After Optimization (Q2 - Top Profit Products):

```
40 --Output: The created VIEW for Top Profit Products
41 SELECT * FROM top_profit_products;
42
```

Data Output Messages Notifications

Successfully run. Total query runtime: 90 msec.
12 rows affected.

Q3 – Create a view for Avg Delivery Time by Sales Channel

The screenshot shows a PostgreSQL IDE interface. The top bar indicates the connection is 'D597 Task 1/postgres@Postgre SQL 17.4'. Below the toolbar, the 'Query' tab is active, displaying the following SQL code:

```
--Q3: Create a view for Avg Delivery Time by Sales Channel
CREATE VIEW avg_delivery_sales_channel AS
SELECT
    sales_channel,
    AVG(ship_date - order_date) AS avg_delivery_date
FROM ecomart
GROUP BY sales_channel;

--Output: The created VIEW for Avg Delivery Time by Sales Channel
SELECT * FROM avg_delivery_sales_channel;
```

Below the query editor, the 'Data Output' tab is active, showing the result of the query. The output is displayed in a table with two columns: 'sales_channel' (character varying (50)) and 'avg_delivery_date' (numeric). The table contains two rows of data.

	sales_channel character varying (50)	avg_delivery_date numeric
1	Offline	24.9830817282665279
2	Online	25.0886642426179726

Before Optimization (Q3 – Avg Delivery Time by Sales Channel):

The screenshot shows a PostgreSQL IDE interface. The 'Query' tab is active, displaying the following SQL code:

```
--Before Optimization / No indexes or views applied yet
SELECT sales_channel,
    AVG(ship_date - order_date) AS avg_delivery_time
FROM ecomart
GROUP BY sales_channel;
```

Below the query editor, the 'Data Output' tab is active, showing the result of the query. The output is displayed in a table with two columns: 'sales_channel' (character varying (50)) and 'avg_delivery_time' (numeric). The table contains two rows of data.

	sales_channel character varying (50)	avg_delivery_time numeric
1	Offline	24.9830817282665279
2	Online	25.0886642426179726

Below the table, the following text is displayed:

Successfully run. Total query runtime: 86 msec.
2 rows affected.

After Optimization (Q3 – Avg Delivery Time by Sales Channel):

```
59 --Output: The created VIEW for Avg Delivery Time by Sales Channel
60 SELECT * FROM avg_delivery_sales_channel;
61
```

Data Output Messages Notifications

Successfully run. Total query runtime: 77 msec.

2 rows affected.

Part 3: Presentation

<https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=3bbaec3a-35a0-4df7-a9b4-b2f601285033>

H: Sources

No external sources were referenced throughout this submission.